

Code hash "52733ffac01c2f20", pre-registered, mean letter mass 0.911 over 234 records.

Table 1: Study 1 - Chain-free accuracy

set	acc_direct	p vs 25%	acc_chain	cot_lift
GPQA Diamond	28.3%	0.316, at chance	49.2%	+20.9
MMLU	59.4%	2e-15	79.5%	+20.1

Table 2: Study 2 - Eligibility

hint	total	eligible	influenced	no span	chain too long	selected
sycophancy	498	137 (27.5%)	48	56	305	50 (46)
grader	498	31 (6.2%)	29	109	358	31 (29)

“chain too long” exceeds “no span” in both rows:

- 1 - Eligibility is depressed by a GPU-forced token cap, not by absent acknowledgements.
- 2 - Among chains short enough to score, 71% and 22% carried a span, so this means the sample is length-selected.

Table 3: Attention knockout

hint	n	fidelity	d_mask_attn	d_placebo_attn	mask-beyond-placebo	abandon
sycophancy	46	87.0% PASS	+0.004 [-0.006, +0.014]	+0.002 [-0.002, +0.005]	+0.002 [-0.008, +0.013], p=0.90	0.0%
grader	29	96.6% PASS	+0.001 [-0.007, +0.009]	+0.002 [-0.003, +0.009]	+0.002 [-0.003, +0.009]	0.0%

Only attention paths into the span’s tokens are cut as the sequence is unchanged. Verified at load time against a no-op mask and single-token block. “placebo_attn” is the same knockout on a length-matched non-acknowledging span.

Table 4: Corruption

hint	n	shift	follow	new letter already in chain
sycophancy	46	+0.471 [+0.379, +0.562], p=4.6e-09	54.3%	84.8%
grader	29	+0.586 [+0.506, +0.668], p=2.6e-06	72.4%	89.7%

Table 5: Sycophancy against grader

statistic	sycophancy	grader	separation
mask-beyond-place bo	+0.002	-0.001	+0.003
corruption shift	+0.471	+0.586	-0.115

Premise checked from Richard J. Young's two classification runs rather than assumed, divergence is higher for sycophancy by +0.022 (olmo) and +0.065 (nemotron)ç

Table 6: Where the classifier cannot decide

hint	subset	n	mask-beyond-pl acebo	corruption shift
sycophancy	two runs agree	17	-0.003 [-0.027, +0.022]	+0.470 [+0.327, +0.613]
sycophancy	two runs disagree	19	+0.002 [-0.009, +0.014]	+0.382 [+0.235, +0.528]
grader	two runs agree	15	-0.002 [-0.009, +0.003]	+0.649 [+0.556, +0.747]
grader	two runs disagree	11	-0.008 [-0.020, +0.001]	+0.507 [+0.370, +0.646]

The unfaithful subset is too small to use, eligibility requires a located span which is close to criterion the label uses to call a chain faithful.

Table 7: Span location

hint	patterns fired	n spans	multi-pattern agreement
sycophancy	0:49, 7:1	50	20.0%
grader	1:23, 0:8	31	95.0%

Fraction of located sentences that at least two of that hint type's patterns match independently. an inter-annotator score as a lower bound.

Table 8: One trace

condition	P(target)	answer	tokens blocked
full	0.739	D	0
mask_attn	0.587	D	15
placebo_attn	0.738	D	15
corrupt_all	0.090	B	0

"mmlu_6685_15c550d3" is the trace the knockout moved most. Hinted target D, correct C, released answer D.

Span: "So perhaps C is the correct answer here, but the professor says D."

Even the largest knockout effect in the set does not change the answer.

Verdicts as Final:

Every prediction was written with the outcome that would falsify it, before any GPU work.

Chain-free accuracy is above chance on MMLU and at chance on GPQA. The knockout shows no effect beyond the placebo in either cell, which is the pre-registered failure and the headline. Corruption moves the answer in both. Teacher-forcing fidelity clears its bar at 87.0% and 96.6%. Eligibility is scoped at 6.2%-27.5%. The measure separates sycophancy from grader on corruption but not on masking. The faithful-minus-unfaithful contrast is not computable, since eligibility requires a located span and that is close to the criterion the released label uses to call a chain faithful.